



# 2016中国开源年会

## China Open Source Conference 2016

时间：2016年10月15日-16日

地点：北京航空航天大学

# 基于DC/OS的现代应用 实践

Sam Chen

2016.10

# 个人简介



- DC/OS中国社区创始人和社区项目孵化者。
- 陈冉(Sam)先生曾任职Linker Networks首席技术官和技术副总裁，主要负责Linker 的技术路线制定，新技术孵化和产品线推广。在这之前，在惠普云事业部担任首席架构师和首席技术官。
- Mesos/OpenStack/Docker/CF社区共同负责人。
- OSCAR DCOS PL/Core。
- 开源布道者。
- Linux 基金会中国协调人

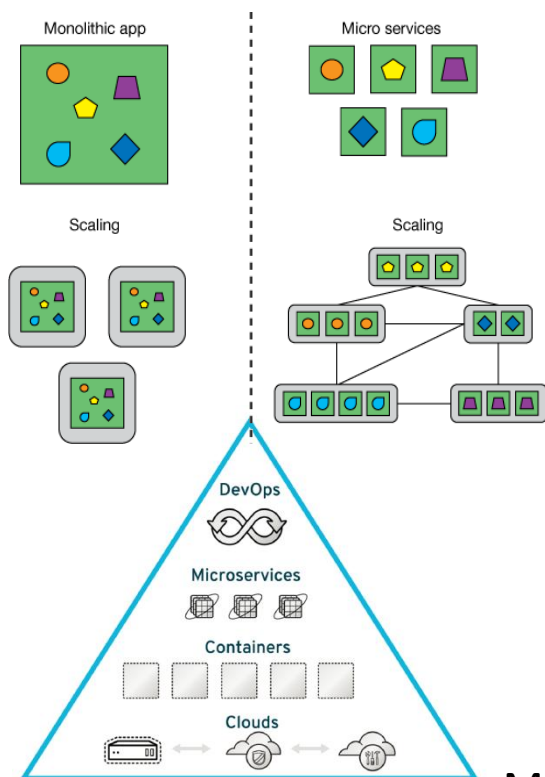




# 演讲内容

- 概念解释
- 微服务
- 现代应用
- DC/OS是什么
- DC/OS vs. Docker vs. K8S vs. Openstack
- 如何现实现代应用在DC/OS架构
- DC/OS为什么 “与众不同”
- 落地最佳实践
- 案例分析

# 概念介绍



## Services & Applications

Easily deploy and run datacenter-wide app services such as Docker, Cassandra, and Spark pooled on a single platform

## Run Anywhere

Bare-metal, virtual, cloud or hybrid  
- DC/OS runs on it all - only requirement is a modern Linux distro; Windows support coming soon



DC/OS Powered by Apache Mesos

Runtime, tools and best practices built-in to simplify operations and deliver a production self-healing infrastructure

Modern Apps= (DC/OS + Micro services + CI/CD + Hybrid Cloud)

# 微服务要素



## The Twelve Factors

### 1. Code Base

One codebase tracked in revision control, many deploys

### 2. Dependences

Explicitly declare and isolate dependences

### 3. Config

Store config in environment

### 4. Backing Services

Treat backing services as attached resources

### 5. Build, release, run

Strictly separate build and run stages

### 6. Process

Execute the app as one or more stateless processes

### 7. Port Binding

Export services via port binding

### 8. Concurrency

Scale out via process model

### 9. Disposability

Maximize robustness with fast startup and graceful shutdown

### 10. Dev/prod parity

Keep development, staging, and production as similar as possible

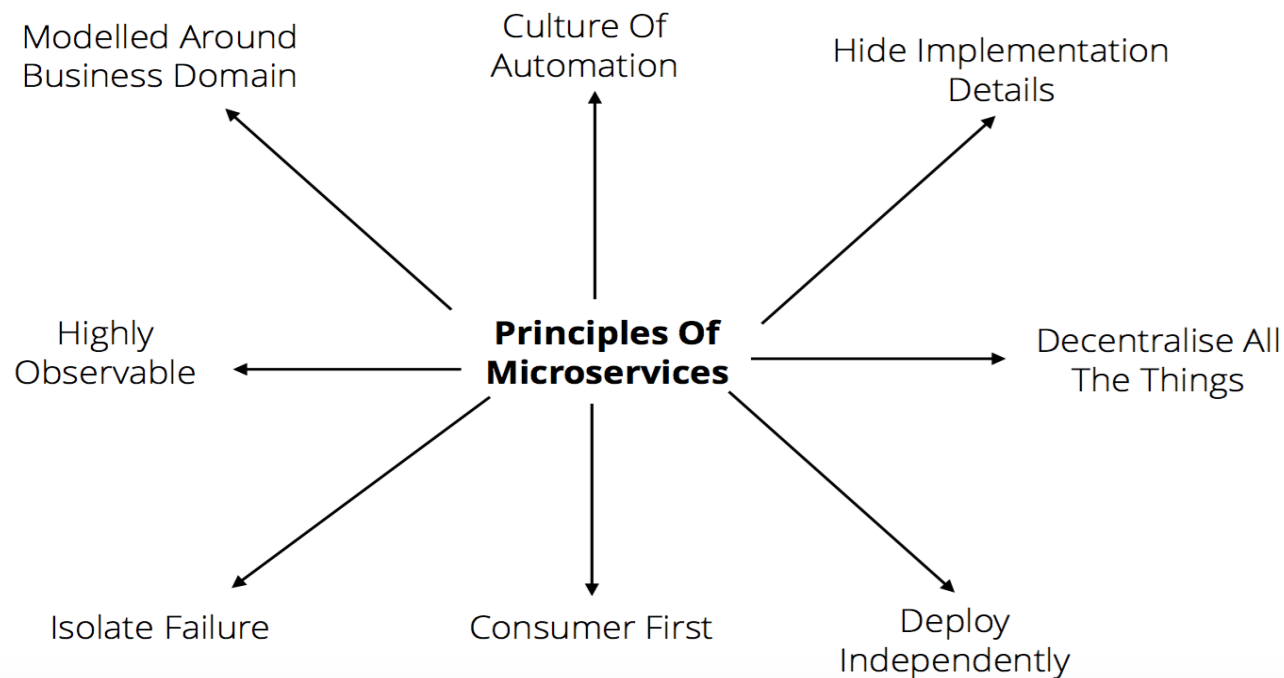
### 11. Logs

Treat logs as event stream

### 12. Admin processes

Run admin/management tasks as one-off processes

# 微服务准则





# 什么是现代应用



UNIT OF  
INTERACTION

NEW FORM FACTOR FOR  
DEVELOPING AND  
RUNNING APPS

PARTITION (LPAR)

SERVER

VIRTUAL MACHINE

**DATACENTER**



**MAINFRAME**

**PHYSICAL (x86)**

**VIRTUAL**

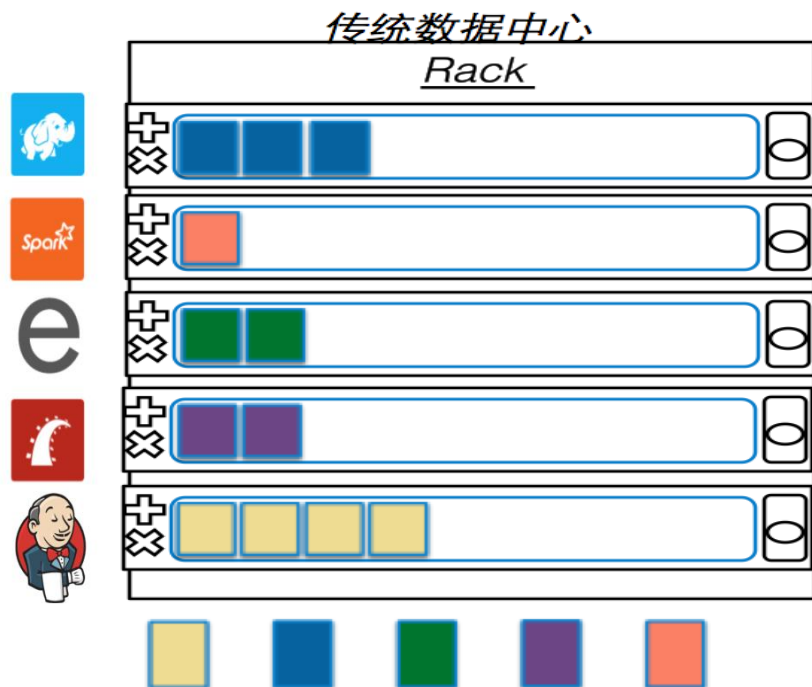
**HYPERSCALE**

- DATA / TRANSACTION PROCESSING
- UNIX, IBM OS/360
- ERP, CRM, PRODUCTIVITY, MAIL & WEB SERVER
- LINUX, WINDOWS
- ERP, CRM, PRODUCTIVITY, MAIL & WEB SERVER
- HYPERVISOR + GUEST OS
- BIG DATA, INTERNET OF THINGS, MOBILE APPS
- **DATACENTER OPERATING SYSTEM (DC/OS)**


DEFINITIVE



# DC/OS是什么



Colorful rectangles are any distributed services running on pre-allocated servers with characteristics of low resource utilization and complicated operation management

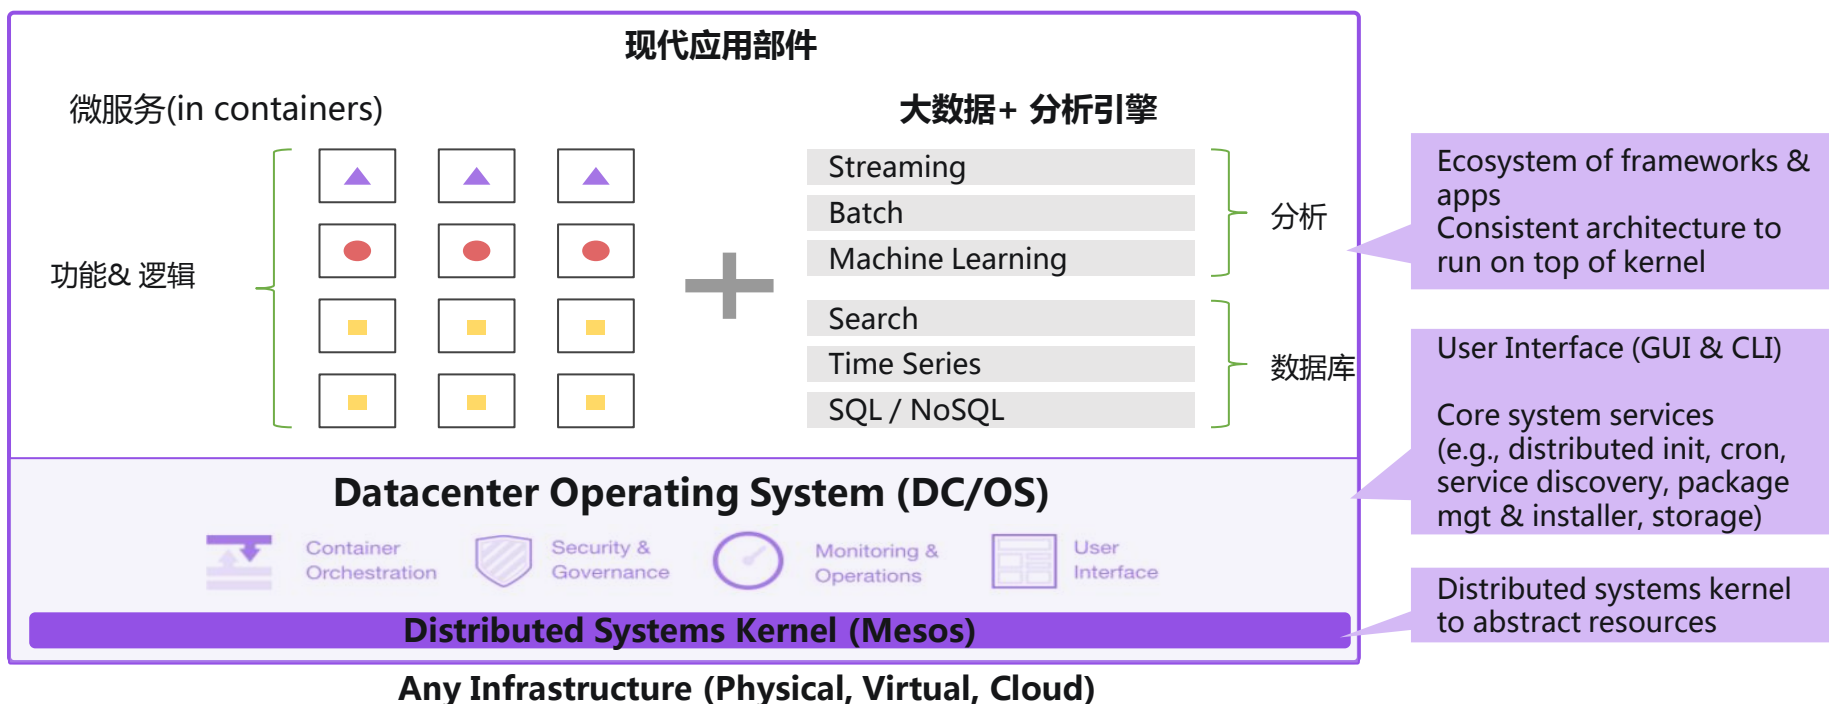


DC/OS maximizing resource utilization with higher availability, elasticity and robust in fabric and aggregated datacenter

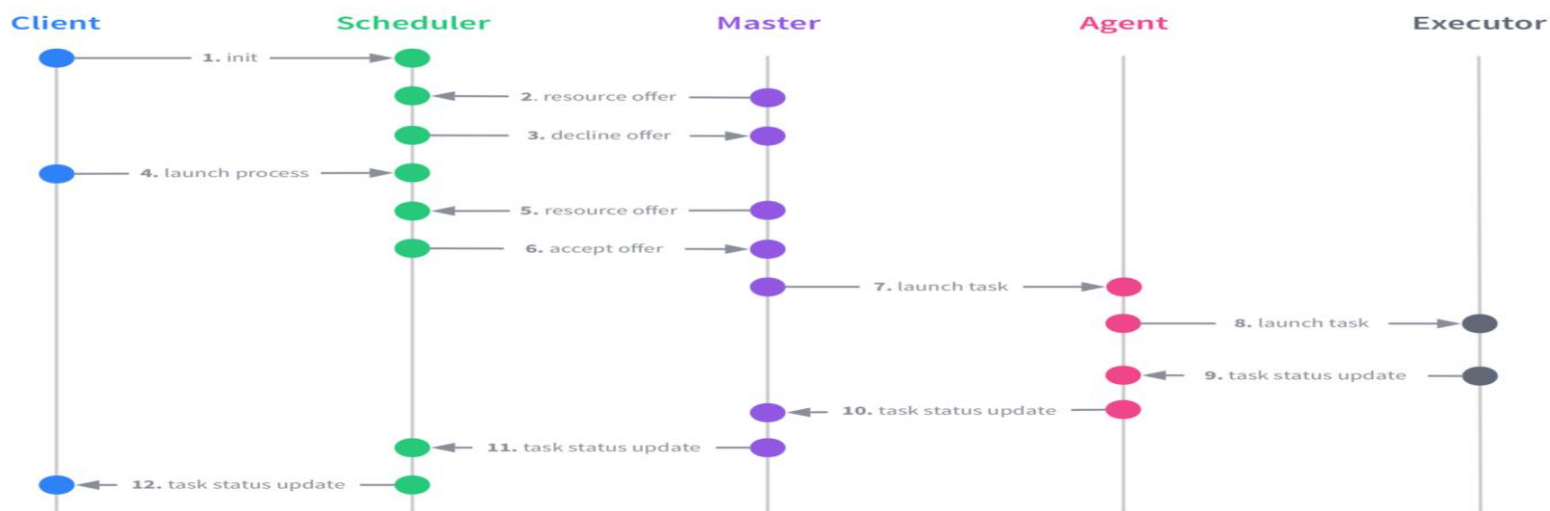
# DC/OS与现代应用



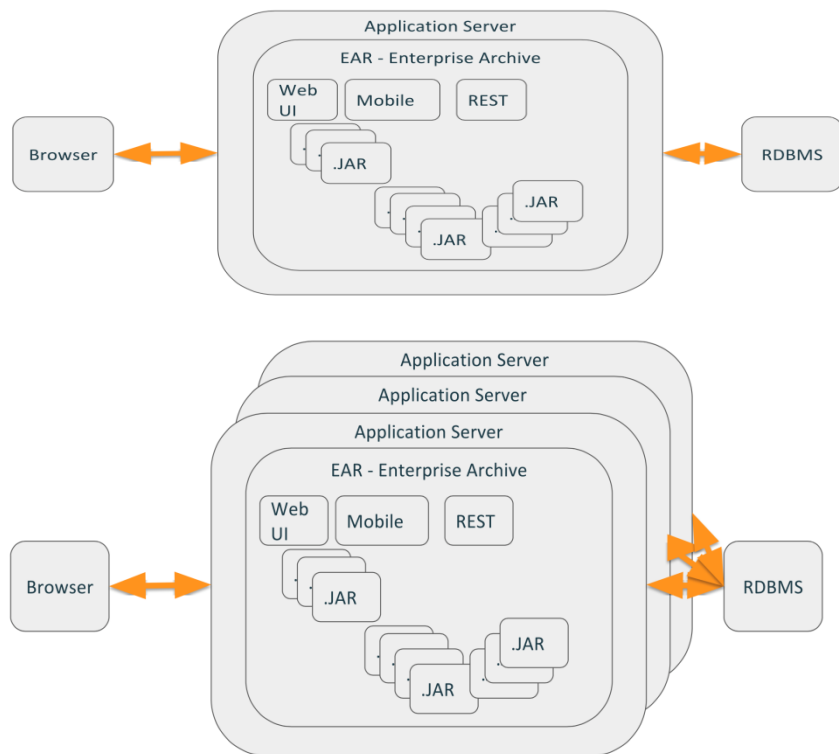
- 100% 开源(ASL2.0) + 一个生态 + 一个多元化的社区
- 在此之下有 ~30 OSS 项目 + 路线图和设计 + The build tool chain + Docs and tutorials + Not limited in a



# DC/OS在现代应用中的优势



# 传统应用现状



- Monolithic application – everything is package into a single .ear
- Reuse primarily by sharing .jars
- A “big” push to production once or twice a year
- Single database schema for the entire application
- $\geq$  Heavyweight Infrastructure
- Thousands of Testcases
- Barely New Testcases
- $\geq 20$  Team Member
- The single .ear requiring a multi-month test cycle /
- Huge bug and feature databases
- User Acceptance Undefined
- Technical Design Approach
- Barely Business Components or Domains
- Requiring multiple team involvement & significant oversight
- Technical Dept
- Grown applications
- Outdated Runtimes (Licenses, Complex Updates)

# 传统应用新需求



- 研发部门期待缩短从开发到产品的生命周期，更好服务公司业务降低风险；
- 激烈竞争的商业市场下，GTM期待传统应用快速演进，能够满足市场需求；
- 前端部门对终端用户的无间断，高可靠性的服务要求不断加强，提供竞争力；
- 在IT和数据中心上的费用不断攀升，期待降低运维和运营成本；
- 传统企业期待业务系统能够根据Data的分析，从而持续修正自己；



# 传统应用 > > > 现代应用

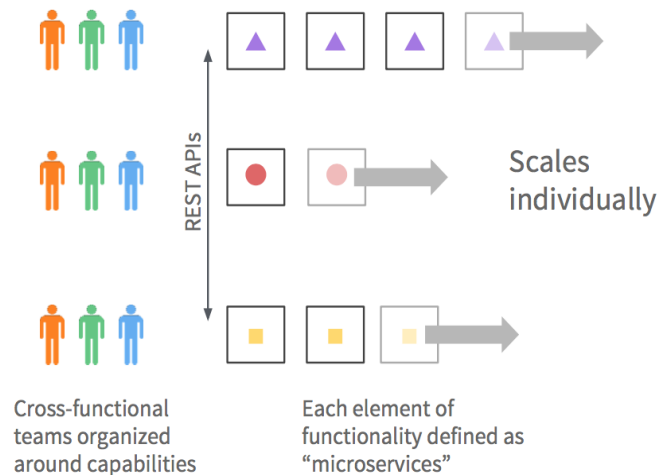


## Traditional Architecture



Small number of large processes with strong inter-dependencies

## Microservices Architecture



Cross-functional teams creating new microservices without interdependencies

# 生产环境验证过的伸缩性



## DC/OS FEATURE

- DC/OS is built on the production-proven Apache Mesos and Marathon
- Built and refined with best practices by some of the world's top experts in distributed systems
- Tested to 10s of thousands of nodes in production for more than 6+ years in production
- Designed for resiliency and scale from day 1



Apache  
**MESOS**™

## BENEFITS

- Reliable, mature technology with years of bug fixing, reduce risk of downtime for mission critical workloads
- Easily expand from small to very large clusters with confidence
- Open source core prevents vendor lock-in



# 现代应用分布式高可用性

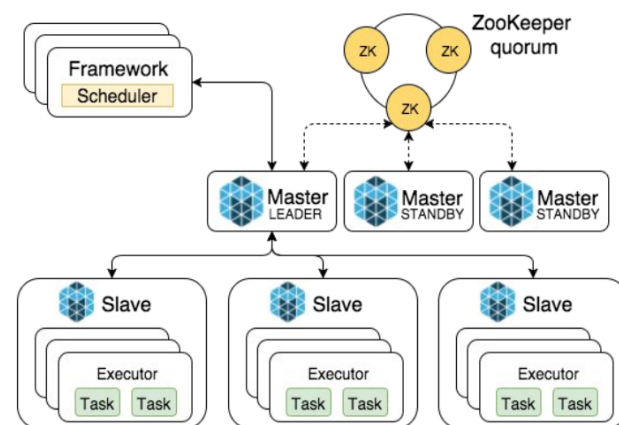


## DC/OS FEATURE

- No SPOF, N+1 Architecture
- Integrated KV Store, DB & load balancing for seamless failover and maintenance
- Turn-key solution, no external dependencies for HA
- Agent crash, restart & upgrade does not affect workloads
- Erlang based distributed networking stack for on the fly upgrade
- Non disruptive failover and upgrade for masters and agents

## BENEFITS

- Highest uptime for scalable masters, agents and workloads
- No manual intervention in case of master/agent failover
- Easily move to the latest DC/OS release
- Resilience to network partitioning & split brain scenarios



# 容器和数据服务



## DC/OS FEATURE

### Supports Stateless and Stateful Containers

- Docker Containers
- Mesos Containers (i.e Java Binaries, Go Binaries)

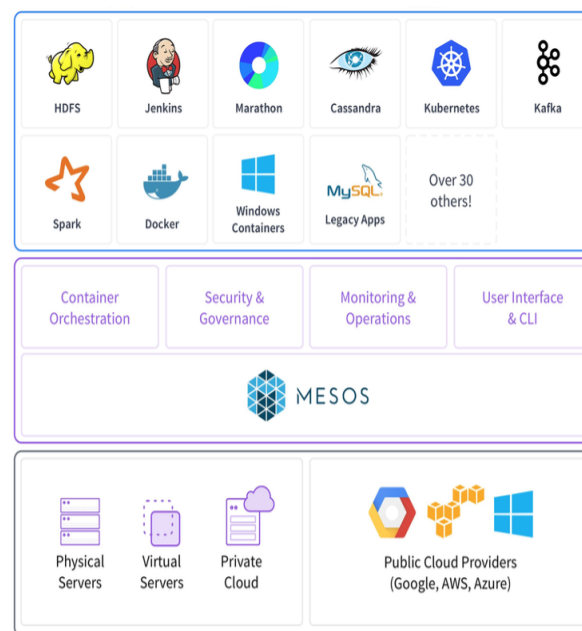
### Supports Stateful Data and Analytics Services

- Analytics & Big Data (i.e Spark, HDFS), No-SQL Databases (i.e Cassandra)
- Search (i.e Elasticsearch), Message Queueing (i.e Kafka)

### Linux and (soon) Windows

## BENEFITS

- Increased Utilization across datacenter
- Easily scale applications up and down
- Single platform for all applications, simplifying operations



# 服务发现

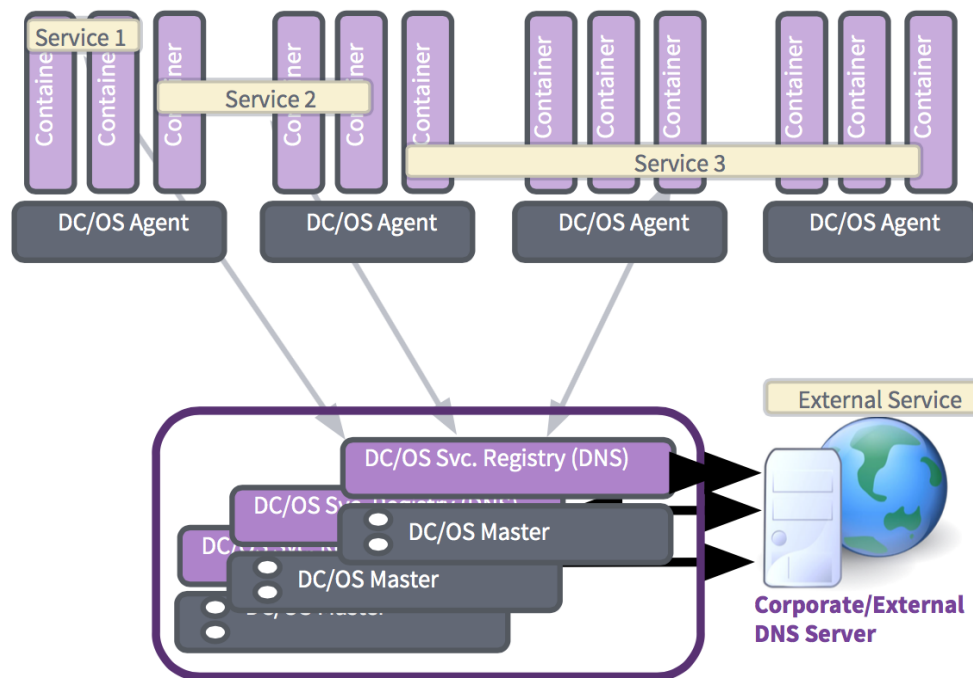


## DC/OS FEATURE

- DNS-based Service Registry & Distr. DNS Proxy
- Integrated DC/OS Service Discovery
- Integrated External Service Discovery

## BENEFITS

- Highly available & scalable service registry
- Turn-key solution to deploy micro-services
- Integration with existing non-DC/OS services to enable brownfield deployments.




# 可靠性 - L4 LB



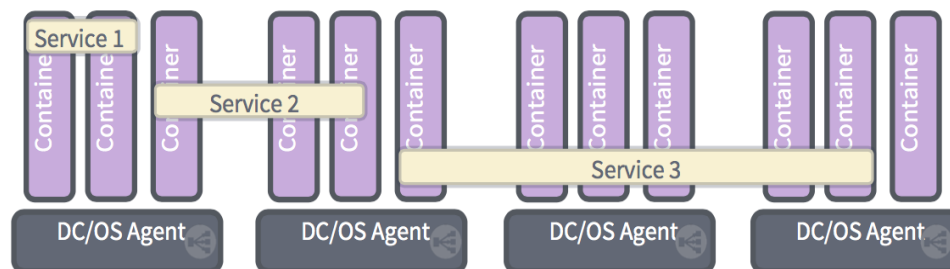
## DC/OS FEATURE

- Fast Converging Distributed Load Balancer
- Integrated with DC/OS Service Discovery
- Integrated Blue-Green deployment
- Support for variety of Layer 4 LB algorithms



## BENEFITS

- Highly available LB with no single choke point
- Enable non-disruptive service upgrades with blue-green deployments.
- Highly scalable and tolerant to large # of host failures.




# 有状态容器和卷




## DC/OS FEATURE

- Persistent Local Volumes for Containers
- External Volumes using Shared Storage
- DC/OS Volume Management UI



## BENEFITS

- Run Stateful applications like MySQL in Containers.
- Integrate with existing shared storage
- Manage & Troubleshoot with integrated volume manager





# “三国杀” + ?



DC/OS



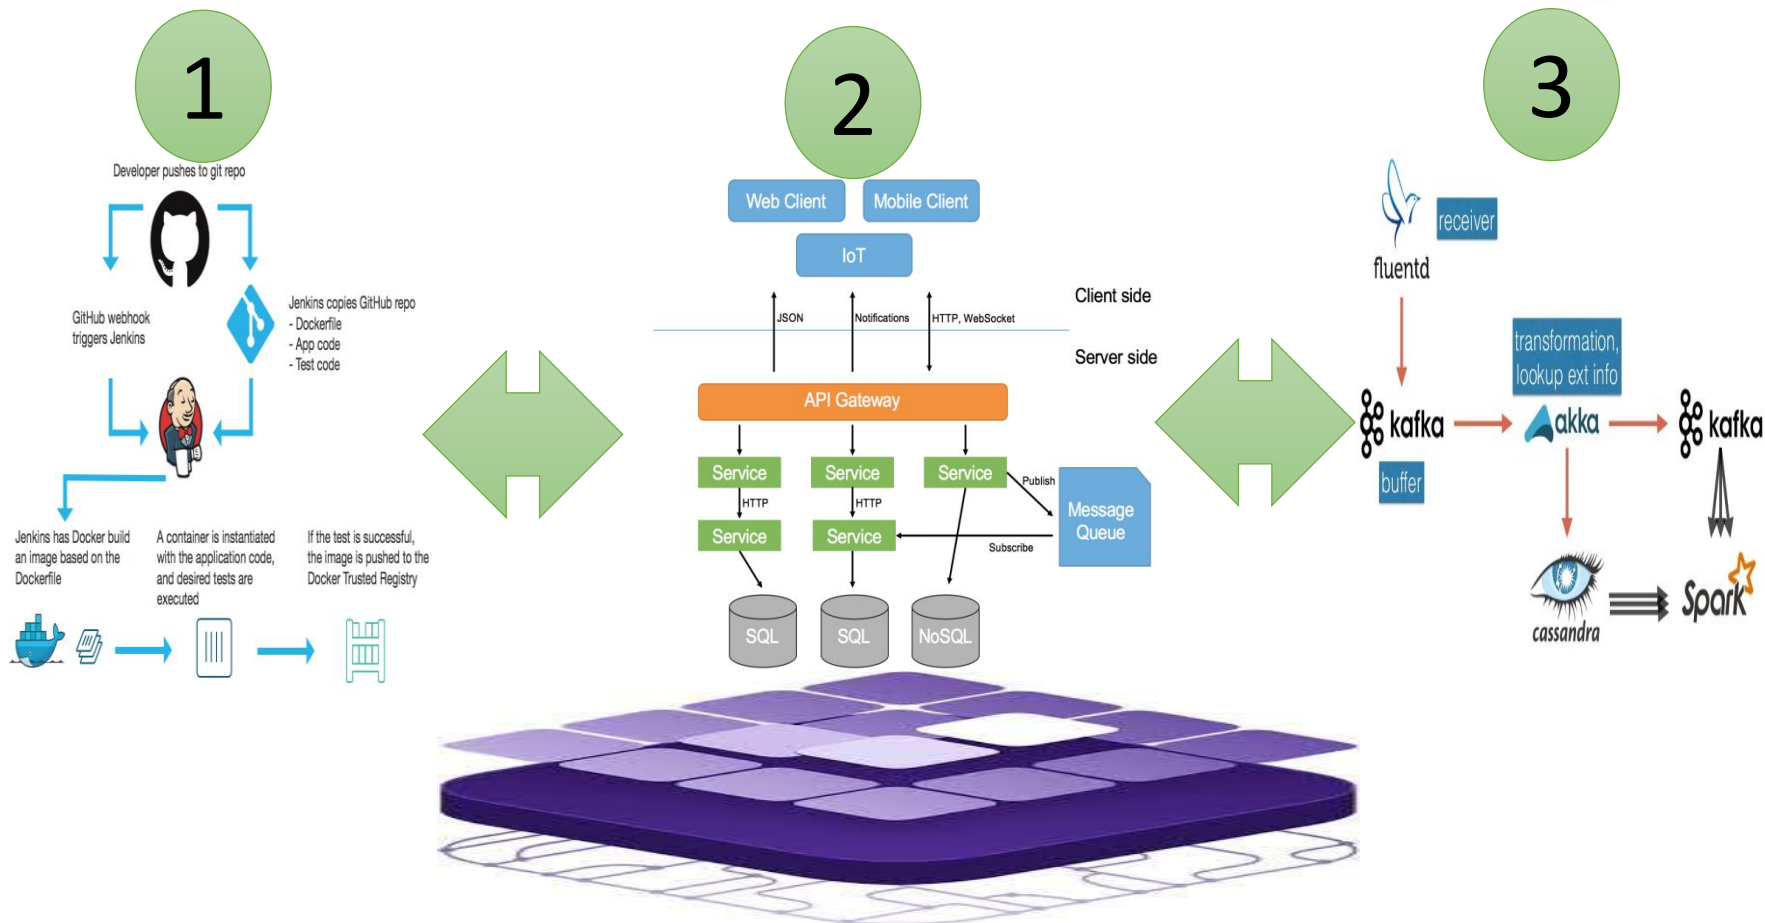
kubernetes



openstack™  
CLOUD SOFTWARE



# 现代架构最佳实践





# 案例分享一



## Challenges

- 400 developers submitting jobs to Jenkins for CI/CD builds experienced sizable delays in task completion due to Jenkins job queuing

## Mesosphere Solution

- Customer wanted to replicate the eBay use case<sup>1</sup> of running CI builds for eBay's applications in Mesos with Docker containers
- Mesosphere DC/OS allowed Customer to move from an enterprise Continuous Integration solution to open source as Marathon provides equivalent HA functionality

## Challenges

- Needed a production grade native container service that would work on premises and on azure, at massive scale
- Must easily integrate with Azure CI/CD, app management and auto scaling infrastructure
- Microsoft and Linux friendly technology

## Mesosphere Solution

- After independent evaluation, MS team determined Mesos/Mesosphere was the right fit
- Currently integrating Mesosphere DC/OS as a core technology for Azure Container Service

- One of North America's leading diversified financial services companies
- Provides banking, wealth management, insurance and capital markets services on a global basis



## Challenges

- Two Sigma OPS struggled with developer demands for agile real-time analysis
- Already explored various IaaS & PaaS solutions

## Mesosphere Solution

- After successful consulting services engagement determined Mesosphere was the right fit
- Compelling reason to move fast; Agility/Performance/Scalability
- 100s of servers moving to 1000s in next 6 months

## Challenges

- Verizon needed infrastructure that could handle the volume and speed of data that its users generate across video services and mobile apps
- Verizon was seeking to improve automation, scalability and efficiency when deploying applications, services and big data infrastructure

## Mesosphere Solution

- Mesosphere DC/OS allowed Verizon to quickly launch new products and services while reducing the IT requirements in their datacenters
- Chose Mesosphere DC/OS for hybrid cloud capabilities, to move from AWS to Verizon's private datacenter



- NYC-based Hedge Fund
- \$25 Billion AUM
- Uses a variety of technological methods for its trading strategies:
  - Artificial Intelligence
  - Machine Learning
  - Distributed Systems



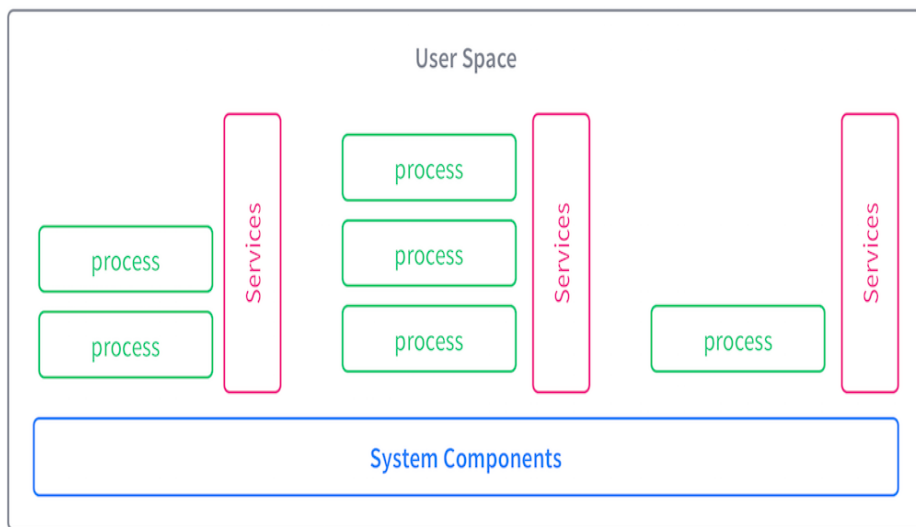
Larry Rau from @Verizon with @flo Launching 50,000 containers in seconds with @mesosphere #DC/OS



# 案例分享二

- 电信运营商
- 银行
- 互联网公司
- 保险
- 车联网

# DC/OS服务集



## System Components are installed and are running by default in the DC/OS cluster and include the following:

The **Admin Router** is an open source NGINX configuration that provides central authentication and proxy to DC/OS services.

**Exhibitor** automatically configures ZooKeeper during installation and provides a usable Web UI to ZooKeeper.

**Mesos-DNS** provides service discovery, allowing apps and services to find each other by using the domain name system (DNS).

**Minuteman** is the internal layer 4 load balancer.

Distributed DNS Proxy is the internal DNS dispatcher.

DC/OS Marathon, the native Marathon instance that is the 'init system' for DC/OS, starts and monitors DC/OS services.

ZooKeeper, a high-performance coordination service that manages the DC/OS services.

The **Cosmos service** is our internal packaging API service.

The diagnostics service (also known as **3DT** or `dcos-ddt.service`, no relationship to the pesticide!) is our diagnostics utility for DC/OS systemd components.

**Logrotate service** does what you think it does: ensures DC/OS services don't blow up cluster hosts with too much log data on disk.

**Marathon** shouldn't need any introduction, it's the distributed init system for the DC/OS cluster.

The DC/OS **signal service** queries the diagnostics service `/system/health/v1/report` endpoint on the leading master and sends this data to SegmentIO for use in tracking metrics and customer support.

The **history service** provides a simple service for storing stateful information about your DC/OS cluster.



# Join the DC/OS Community

Connect with our community of users and browse the latest DC/OS news.



## GitHub

Are you interested in helping us make DC/OS even better? Let's work together! Check out our source code on GitHub.

**[View repositories →](#)**



## Slack

Have any questions? Our Slack channel is the best place to get help. Just send us a request to automatically receive your invitation.

**[Join chat →](#)**



## Mailing List

Want to stay in the loop and connect with other community members? Our public mailing list has all the latest updates. Join the discussion.

**[Join users@dcos.io →](mailto:users@dcos.io)**

# DCOS.io



The screenshot displays the DCOS.io website interface. At the top, the navigation bar includes links for 'Why DC/OS', 'Install', 'Docs', 'Community', 'GitHub', and a 'Get Started' button. The main headline reads 'The easiest way to run containers in production.' with a 'Get Started' button and a 'Play Video' link. Below this, a grid of service cards is shown, each with an icon, name, and brief description:

- Marathon**: Use to launch long-running apps, for example a webserver.
- Spark**: Use for elastic data processing, batch or streaming.
- ArangoDB**: Use to store and query tree-like data as well as graphs.
- Cassandra**: Use to store and query semi-structured data in wide tables.
- Chronos**: Use to launch scheduled batch jobs in a distributed way.
- Jenkins**: Use to build your scalable and fault-tolerant CI/CD pipeline.
- Kafka**: Use to queue messages and distributed to others.

Below the service cards, four key features are highlighted:

- Containers & Big Data**: Easily deploy and run stateful or stateless distributed workloads including Docker containers, big data, and traditional apps.
- 100% Open Source**: Run DC/OS in production without any scale or performance limitations, on any on-premises or cloud platform.
- Production Proven**: Built on Apache Mesos with the experience of Mesosphere, Yelp, Twitter, Airbnb, and many of today's most innovative companies.
- Datacenter Apps**: Build apps quickly using a rich ecosystem of datacenter-scale application services.



# 谢谢



Sam

Dongcheng District, Beijing



Scan the QR code to add me on WeChat

## Q&A